

# Efficient Attention Mechanisms: BigBird and FlashAttention

Mohamed Ramadan

## 1 Introduction

Transformer models rely on the self-attention mechanism, where each token attends to every other token in the sequence.

Given a sequence of length  $n$ , full self-attention computes:

$$\text{Attention}(Q, K, V) = \text{Softmax} \left( \frac{QK^T}{\sqrt{d}} \right) V$$

This requires computing an  $n \times n$  attention matrix.

**Problem:**

- Time Complexity:  $\mathcal{O}(n^2)$
- Memory Complexity:  $\mathcal{O}(n^2)$

This becomes infeasible for long sequences (e.g., documents, DNA, video frames).

## 2 Full Attention Bottleneck

In full attention, every token interacts with every other token.




Figure 1: Full Self-Attention: Each token attends to all tokens ( $\mathcal{O}(n^2)$ ).

For example, with 10 tokens:

$$10 \times 10 = 100 \text{ attention scores}$$

For 10,000 tokens:

$$10^4 \times 10^4 = 10^8$$

Clearly quadratic growth is the bottleneck.

### 3 BigBird: Sparse Attention for Linear Scaling

BigBird introduces structured sparse attention to reduce complexity from  $\mathcal{O}(n^2)$  to  $\mathcal{O}(n)$ .

Instead of attending to all tokens, each token attends to:

- **Global tokens**
- **Local window tokens**
- **Random tokens**

$$\text{Complexity} = \mathcal{O}(n)$$

#### 3.1 Intuition

- Global tokens maintain long-range communication.
- Local window captures nearby context.
- Random connections ensure theoretical expressivity.




Figure 2: BigBird Sparse Attention Pattern (Global + Local + Random).

## 3.2 Why BigBird Works

BigBird was proven to be:

- A universal approximator of sequence functions
- Turing complete under certain conditions

Thus, it preserves theoretical power while reducing computational cost.

## 4 FlashAttention: IO-Aware Exact Attention

FlashAttention does **NOT** approximate attention.

Instead, it computes exact attention but optimizes memory access.

### 4.1 Core Idea

The real bottleneck is memory bandwidth (HBM), not FLOPs.

FlashAttention:

1. Tiles the attention matrix into blocks
2. Computes each block inside fast SRAM
3. Uses an online softmax trick to merge results
4. Avoids materializing the full  $n \times n$  matrix

Time Complexity:  $\mathcal{O}(n^2)$

Memory Usage: Dramatically Reduced




Figure 3: FlashAttention: Tiled computation inside GPU SRAM.

## 4.2 Key Insight: Online Softmax

Instead of computing:

$$\text{Softmax}(QK^T)$$

FlashAttention maintains a running max and running sum across tiles, ensuring numerical stability and exact equivalence to full attention.

Thus:

**Same result as full attention, faster and memory efficient**

## 5 BigBird vs FlashAttention

	BigBird	FlashAttention
Attention Type	Sparse	Exact
Time Complexity	$\mathcal{O}(n)$	$\mathcal{O}(n^2)$
Memory Efficiency	High	Very High
Long Sequence Scaling	Excellent	Limited by quadratic FLOPs
Accuracy vs Full	Approximate	Identical
Best Use Case	Very long documents	Large LLM training/inference

Table 1: Comparison between BigBird and FlashAttention

## 6 When to Use Each?

**Use BigBird When:**

- Working with extremely long sequences (8k–100k tokens)
- You need linear scaling
- Small approximation is acceptable

**Use FlashAttention When:**

- Training large LLMs
- You want exact attention
- GPU memory bandwidth is bottleneck

## 7 Real-World Applications

**BigBird Used In:**

- Long document classification
- Question answering over long contexts
- Genomics

**FlashAttention Used In:**

- GPT-style models

- LLaMA variants
- Modern production LLM systems

## 8 Conclusion

Full attention is powerful but quadratic.

BigBird solves the scaling problem via sparse structure.

FlashAttention solves the memory bottleneck via IO-aware optimization.

They address different dimensions of the same core challenge.

**BigBird = Algorithmic Efficiency**

**FlashAttention = Hardware Efficiency**